

06-12-2023-2

Hengde Ouyang

2023-06-12

2. Horseshoe Prior Investigation

$$\beta_j | \lambda_j, \tau \sim N(0, \tau^2 \lambda_j^2)$$

$$\lambda_j \sim C^+(0, 1), \quad j = 1, \dots, p$$

In our function, we define:

$$\sigma_j = \tau \lambda_j$$

1. Horseshoe prior is well-known with its shrinkage ability. In our simulation, it indeed shows this ability.

Function:

```
library(LaplacesDemon)

##
## Attaching package: 'LaplacesDemon'

## The following objects are masked from 'package:mvtnorm':
##
##      dmvt, rmvt

MH_horseshoe_Sampling <- function(Y,delta,tau,
                                  A,beta0,sigma0,var.prop,
                                  m,B,eta,
                                  Wmat_option=0){

  accept_beta = 0
  accept_lambda = 0
  beta = beta0
  lambda = lambda0
  sigma = sigma0

  # What we want to record
  BETA = matrix(0,m,dim(A)[2])
  LAMBDA = matrix(0,m,dim(A)[2])
  ThetaRecord <- matrix(0, m, length(Y))
```

```

C_stat = c()

# For safety m>B
if (B>m){
  B = 0
}

# 0 means we use Harrell C statistics
# 1 means we use Uno C statistics
if (Wmat_option==0){
  Wmat <- HarrellC_Wmat(Y, delta, tau)
}else if (Wmat_option==1){
  Wmat <- UnoC_Wmat(Y, delta, tau)
}else{ # Other Possible C index...
  Wmat <- HarrellC_Wmat(Y, delta, tau)
}

for (i in 1:m){

  # Sample beta from proposal distribution
  beta.p = t(rmvnorm(1,beta,var.prop))

  # Compute theta from current and last iteration
  theta.p = THETA(A,beta.p)
  theta = THETA(A,beta)

  # Record theta from last iteration
  ThetaRecord[i,] <- theta

  # Compute C-statistics from current and last iteration
  HC.p = HarrellC(theta.p, Wmat)
  HC = HarrellC(theta, Wmat)

  # Record C-statistics from last iteration
  C_stat = c(C_stat,HC)

  lrMH = eta*log(HC.p) +
    sum(dnorm(beta.p,beta0,sigma,log=T))-
    eta*log(HC) -
    sum(dnorm(beta,beta0,sigma,log=T))

  if (log(runif(1))<lrMH){
    beta = beta.p
    accept_beta = accept_beta + 1
  }
}

```

```

BETA[i,] = beta

#####

lambda.p = exp(t(rnorm(dim(A)[2],log(lambda),rep(1,dim(A)[2]))))
sigma.p = lambda.p*beta_tau

lrMH_lambda = sum(dnorm(beta,beta0,sigma.p,log=T))+
               sum(dhalfcauchy(lambda.p,lambda_scale,log = T))-
               sum(dnorm(beta,beta0,sigma,log=T))-
               sum(dhalfcauchy(lambda,lambda_scale,log = T))

if (log(runif(1))<lrMH_lambda){
  lambda = lambda.p
  sigma = sigma.p
  accept_lambda = accept_lambda + 1
}
LAMBDA[i,] = lambda
}

#####

if (B == 0){
  return(list(BETA=BETA,
              LAMBDA = LAMBDA,
              accept_beta=accept_beta/m,
              accept_lambda=accept_lambda/m,
              THETA = ThetaRecord,
              C_stat = C_stat))
}else{
  return(list(BETA=BETA[-c(1:B)],,
              LAMBDA = LAMBDA[-c(1:B)],,
              accept_beta=accept_beta/m,
              accept_lambda=accept_lambda/m,
              THETA = ThetaRecord[-c(1:B)],,
              C_stat = C_stat[-c(1:B)]))
}

}

system.time({
result_horseshoe = MH_horseshoe_Sampling(Y,delta,tau,
                                         A,beta0,sigma0,var.prop,
                                         m,B,eta,
                                         Wmat_option)
result_no_horseshoe = MH_Sampling(Y,delta,tau,

```

```

        A,beta0,sigma0,var.prop,
        m,B,eta,
        Wmat_option)
})

```

```

##      user  system elapsed
## 220.24   22.66   324.18

```

Comparing the posterior mean for each beta w/wo horseshoe prior:

```

posterior_mean_beta = rbind(colMeans(result_no_horseshoe$BETA),
                             colMeans(result_horseshoe$BETA))
colnames(posterior_mean_beta) = c("beta1", "beta2", "beta3", "beta4",
                                   "beta5", "beta6", "beta7", "beta8")
rownames(posterior_mean_beta) = c("w/o hs", "w hs")
posterior_mean_beta

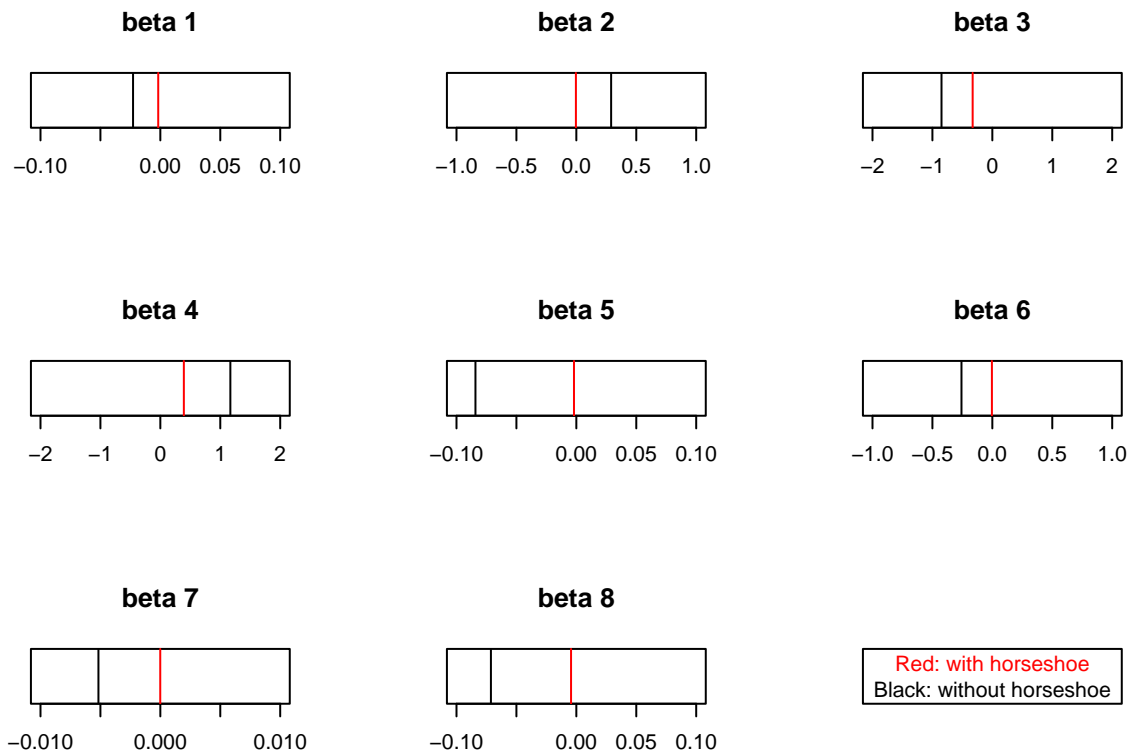
```

```

##              beta1      beta2      beta3      beta4      beta5      beta6
## w/o hs -0.022775302  0.290246059 -0.8485934  1.1683106 -0.084179789 -0.257475138
## w hs   -0.001763546 -0.003673065 -0.3286724  0.3924602 -0.001962776 -0.003677703
##              beta7      beta8
## w/o hs -5.164651e-03 -0.071277654
## w hs   -7.756313e-06 -0.004408845

```

Try to visualize it:



And we can compare the C statistics:

```
Wmat = HarrellC_Wmat(Y,delta,tau)
HarrellC(colMeans(result_no_horseshoe$THETA),Wmat)
```

```
## [1] 0.6426407
```

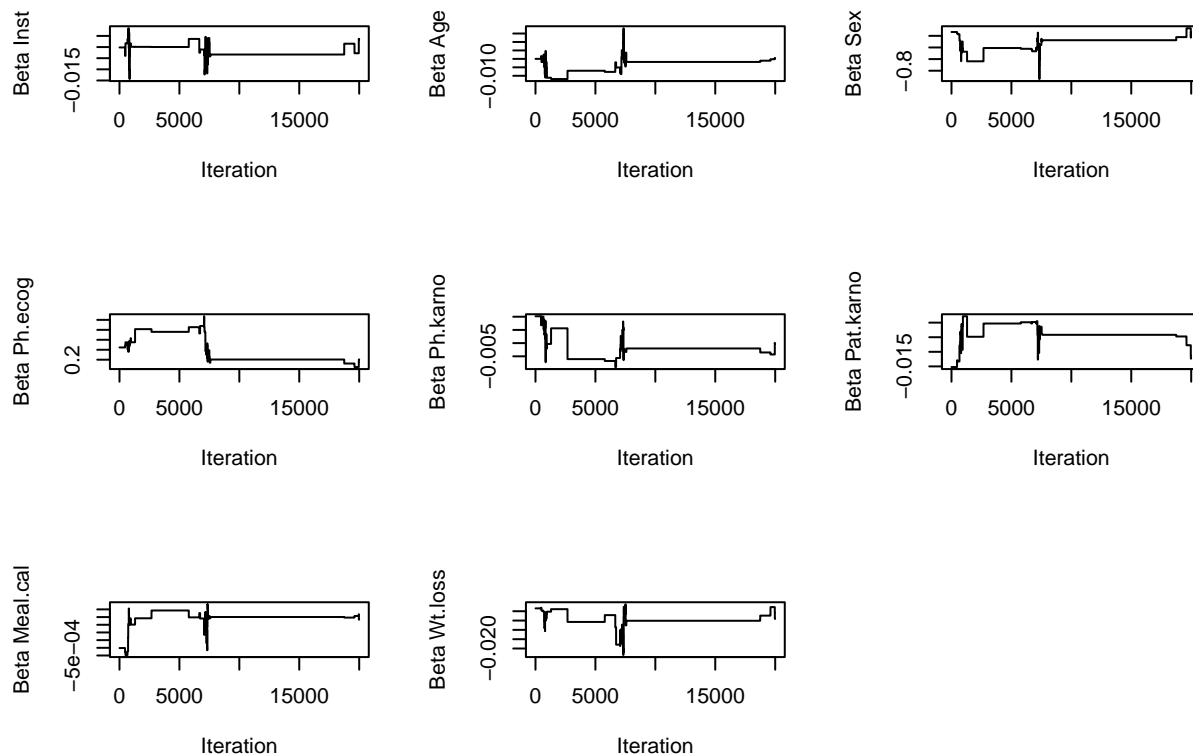
```
HarrellC(colMeans(result_horseshoe$THETA),Wmat)
```

```
## [1] 0.6596893
```

2. However, the issue of a low acceptance rate still remains..

No matter how I tune the parameters, it's still difficult to raise the acceptance rate.

```
beta_iter(result_horseshoe,m,B)
```



Numerical Reason:

The algorithm seems to prefer sampling a low value of lambda. A low value of lambda leads to a small prior variance of beta, which leads to a small acceptance rate. (Based on the simulation result before, smaller prior variance will have smaller acceptance rate)

What's more, when you want to increase beta_tau, the sampling value of lambda will decrease. (No change

in prior variance!)

I also personally searched some papers, trying to find out the reason.

It seems that for some algorithm, this issue will occur when using horseshoe prior: “become stuck after a few thousand iterations”

<https://dl.acm.org/doi/pdf/10.5555/3455716.3455789> (Page 10)

3. I also personally try another approach.

(It might not be valid, but I just want to explore)

Metropolis-Gibbs Algorithm?

```
MH_horseshoe_Sampling2 <- function(Y,delta,tau,
                                   A,beta0,sigma0,var.prop,
                                   m,B,eta,
                                   Wmat_option=0){

  accept_beta = 0
  accept_lambda = 0
  beta = beta0
  lambda = lambda0
  sigma = sigma0

  # What we want to record
  BETA = matrix(0,m,dim(A)[2])
  LAMBDA = matrix(0,m,dim(A)[2])
  V = matrix(0,m,dim(A)[2])
  ThetaRecord <- matrix(0, m, length(Y))
  C_stat = c()

  # For safety m>B
  if (B>m){
    B = 0
  }

  # 0 means we use Harrell C statistics
  # 1 means we use Uno C statistics
  if (Wmat_option==0){
    Wmat <- HarrellC_Wmat(Y, delta, tau)
  }else if (Wmat_option==1){
    Wmat <- UnoC_Wmat(Y, delta, tau)
  }else{ # Other Possible C index...
    Wmat <- HarrellC_Wmat(Y, delta, tau)
  }

  for (i in 1:m){

    # Sample beta from proposal distribution
    beta.p = t(rmvnorm(1,beta,var.prop))
```

```

# Compute theta from current and last iteration
theta.p = THETA(A,beta.p)
theta = THETA(A,beta)

# Record theta from last iteration
ThetaRecord[i,] <- theta

# Compute C-statistics from current and last iteration
HC.p = HarrellC(theta.p, Wmat)
HC = HarrellC(theta, Wmat)

# Record C-statistics from last iteration
C_stat = c(C_stat,HC)

lrMH = eta*log(HC.p) +
      sum(dnorm(beta.p,beta0,sigma,log=T))-
      eta*log(HC) -
      sum(dnorm(beta,beta0,sigma,log=T))

if (log(runif(1))<lrMH){
  beta = beta.p
  accept_beta = accept_beta + 1
}
BETA[i,] = beta

#####
##### For sampling lambda, I try Gibbs sampling this time #####
lambda = sqrt(1/rgamma(dim(A)[2],shape=1,rate=(1/v)+(beta^2/(2*beta_tau^2))))
v = 1/rgamma(dim(A)[2],shape=1,rate=1+1/lambda^2)
sigma = lambda*beta_tau

LAMBDA[i,] = lambda
V[i,] = v
}

#####

if (B == 0){
  return(list(BETA=BETA,
             LAMBDA = LAMBDA,
             V = V,
             accept_beta=accept_beta/m,
             THETA = ThetaRecord,
             C_stat = C_stat))
}else{
  return(list(BETA=BETA[-c(1:B)],

```

```

        LAMBDA = LAMBDA[-c(1:B),],
        V = V[-c(1:B)],
        accept_beta=accept_beta/m,
        THETA = ThetaRecord[-c(1:B),],
        C_stat = C_stat[-c(1:B)])
    }
}

```

This time, I can successfully tune the parameter beta_tau:

```

system.time({
  m = 22000
  B = 2000
  v = rep(1,dim(A)[2])
  beta_tau = 10
  result_horseshoe21 = MH_horseshoe_Sampling2(Y,delta,tau,
      A,beta0,sigma0,var.prop,
      m,B,eta,
      Wmat_option)

  beta_tau = 1
  result_horseshoe22 = MH_horseshoe_Sampling2(Y,delta,tau,
      A,beta0,sigma0,var.prop,
      m,B,eta,
      Wmat_option)

  beta_tau = 0.1
  result_horseshoe23 = MH_horseshoe_Sampling2(Y,delta,tau,
      A,beta0,sigma0,var.prop,
      m,B,eta,
      Wmat_option)

  beta_tau = 0.01
  result_horseshoe24 = MH_horseshoe_Sampling2(Y,delta,tau,
      A,beta0,sigma0,var.prop,
      m,B,eta,
      Wmat_option)

})

```

```

##      user  system elapsed
## 445.91   38.99   683.00

```

```

hs = data.frame(Beta_Tau = c(10,1,0.1,0.01),
  AcceptanceRate=c(result_horseshoe21$accept_beta,
    result_horseshoe22$accept_beta,
    result_horseshoe23$accept_beta,
    result_horseshoe24$accept_beta),
  WOhorseshoe_Beta1 = c(rep(colMeans(result_horseshoe$BETA)[1],4)),
  Whorseshoe_Beta1 = c(colMeans(result_horseshoe21$BETA)[1],
    colMeans(result_horseshoe22$BETA)[1],
    colMeans(result_horseshoe23$BETA)[1],
    colMeans(result_horseshoe24$BETA)[1]))
hs

```


##	Beta_Tau	AcceptanceRate	WOhorseshoe_Beta1	Whorseshoe_Beta1
## 1	10.00	0.86586364	-0.001763546	-0.089500857
## 2	1.00	0.77836364	-0.001763546	-0.070091333
## 3	0.10	0.38918182	-0.001763546	-0.003590347
## 4	0.01	0.02136364	-0.001763546	-0.001424735

We use beta 1 as an example. It seems that acceptance rate is related to the shrinkage effect.
(With a lower acceptance rate, the beta will shrink more close to 0?)